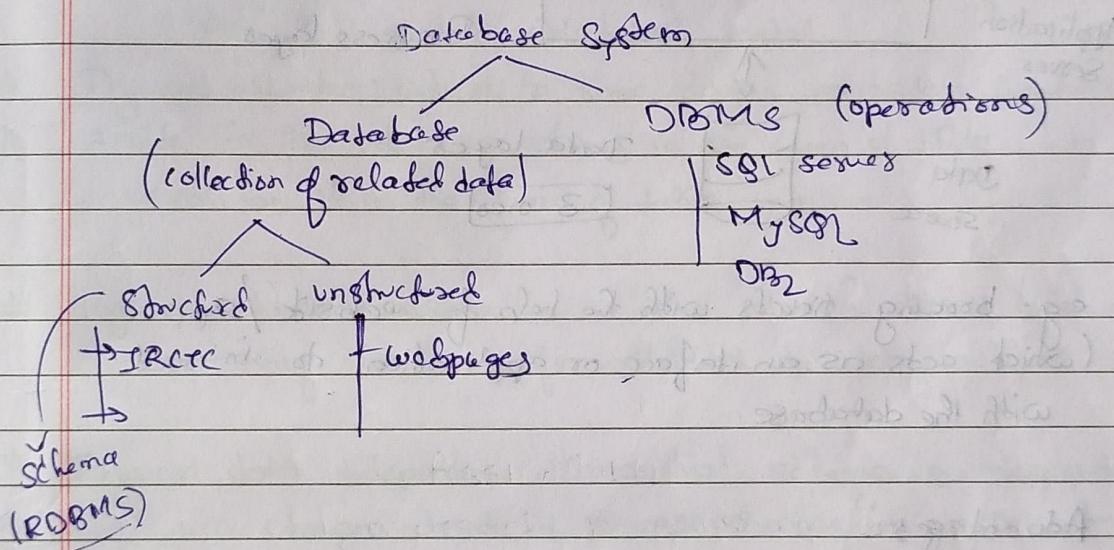


DBMS - Gate 8 summary

- 3 schema, 3-level of abstraction
- Models
- ER Models, Relationships, Keys
- Normalization →
- Transaction control & concurrency (ACID properties)
- SQL & Relational algebra
- Indexing.
 - DDL, DML, DCL, TCL
 - Constraints
 - Aggregate func.
 - Joins
 - Nested query (in, Not in, Any, all)



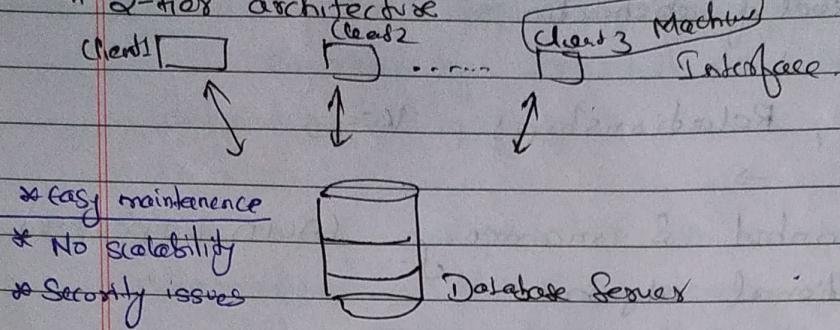
* Structured data is stored in the form of relation.

RDBMS is used to operate on Structured data.

File Systems

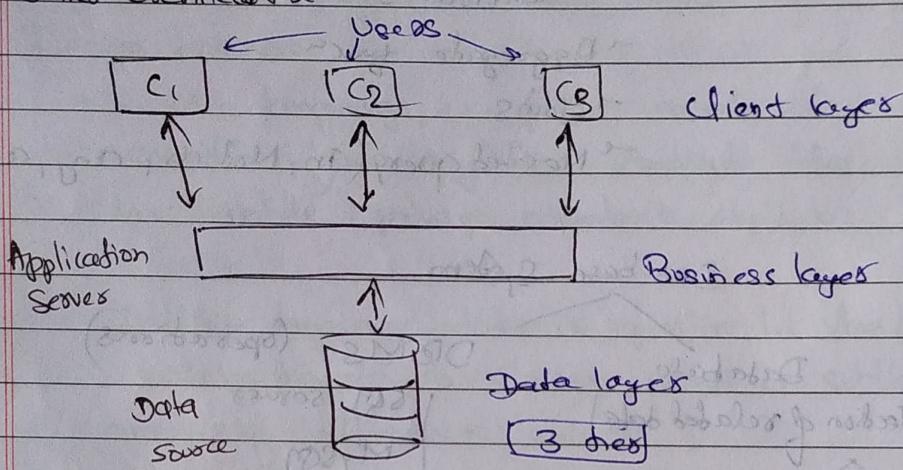
1. In dbms, fast searching is there
2. " " " data is independent of metadata
3. Concurrency: Multiple users can access the db at the same time. Still inconsistency is there with the help of some protocols.
4. Security: (Role based Security)
 - User
 - Admin
 - DBA
 Role based access control.
5. Data redundancy is not there in dbms.

2-tier architecture



e.g. Collecting tickets from window;
Bank transaction

3-tier architecture



e.g. booking tickets with the help of website/app
(which acts as an interface or application) to interact with the database.

Advantages

- Scalable (huge amount of users can be entertained)
- Security (data is not vulnerable to users)

Business layer type processed query, understandable by db.

* 3-tier arc. needs maintenance.

Schemas

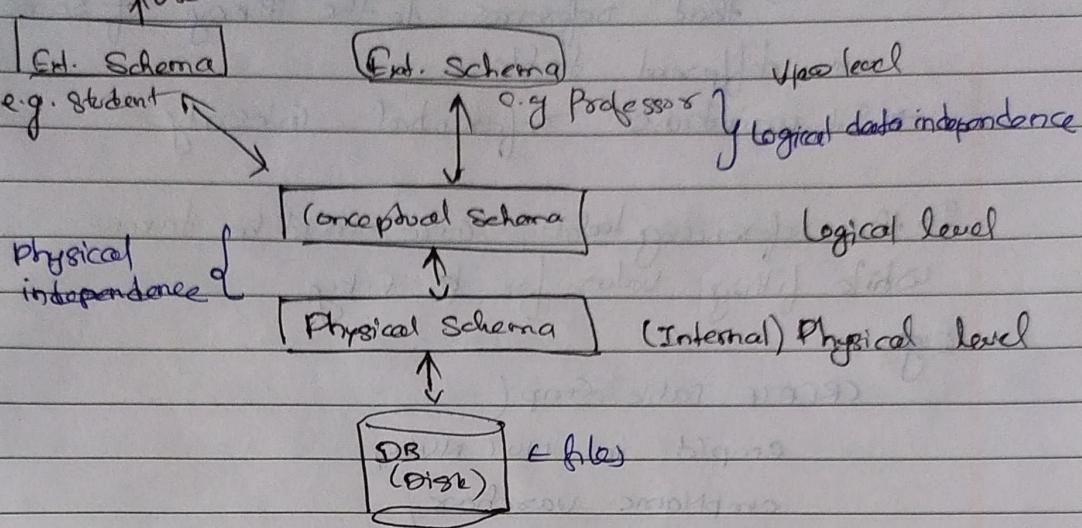
Logical representation of RDBMS

→ Tables

In E-R

→ Entities and Relationships

3-Schema architecture



Conceptual Schema: Structure of representation of data

e.g. E-R model, Relational model.

Physical Schema: Where to store data; particular locations, drives

Inside the DB, data is stored in the form of files, but we can see it in the form of tables etc.

Logical data independence is achieved through views (virtual table having user specific attributes in view table).

Physical data independence: Conceptual schema is independent of changes made in storage structure/indices/Data structure change.

Concept of candidate key

Key → uniquely identify each tuple in a table.
 < set of all such attributes (keys) > is called candidate key.

of superkey

Primary key: smallest subset, which can uniquely identify each tuple.

Alternate keys: All other keys other than P.Key

P.Keys {unique + NOT NULL}

Only one pkey exists for a given relation

Foreign key is an attribute or a set of attributes that references to the P. key of same or different table.

→ It maintains referential integrity

In the referencing table, care must be taken while filling values for P. key.

e.g.

CREATE Table Emp (

empId int NOT NULL,

empName varchar,

PersonId int,

Primary Key (EmpId),

Foreign Key (PersonId) References Persons(PersonId);

If table has been created

ALTER Table Emp

ADD constraint fk

Foreign key (PersonId) References Persons(PersonId);

Referenced table

1. Insert (No violations)

2. Update

3. Delete (Delete all such tuples related to this fkey from referencing table, if any)

Soln:

→ ON delete cascade (delete here as well)

→ ON delete set NULL (if fkey is not primary)

→ ON delete NO Action (in referencing table)

2. Update: (may cause violations)

→ On update

#

Referencing table

1. Insert (May cause violations)

2. Deletion (No issues)

3. Update: Violations ex here

- Superset of candidate key is a superkey
- Primary key is minimal superkey
- Roll no. → Primary key / cand. key
- Roll no. + Name → Superkey.

Table having (A_1, A_2, \dots, A_n) attributes

A_1 → is candidate key.

How many superkey is possible?

2^n (if all are to be chosen or not chosen)

∴ A_1 is key it is to be chosen any how

$$\text{ans} = 2^{n-1}$$

⇒ If A_1, A_2 is key.

$$2^{n-1} + 2^{n-2} = 2^{n-2}$$

$A_1 \quad A_2$ eliminating redundant keys.

ER Model:

→ Entity is any object having physical existence having some attributes.

Relationships: Association among entities.

Type of attributes

1. Single vs Multivalued (RegId vs Phone No)

2. Simple vs Composite (Age vs Address → Lane district)

3. Stored vs Derived (DOB vs Age) → from DOB

4. Key vs Non-key attr.

RegId vs
unique

5. Required vs Optional

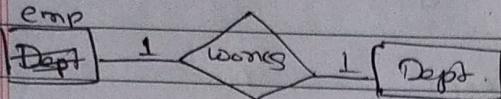
6. Complex

(composite + Multivalued)

e.g. Multiple addresses

Types of relationships

(i) One to One relationships



If must has Pkey of all participating entities + [Descriptive attributes] (optional)

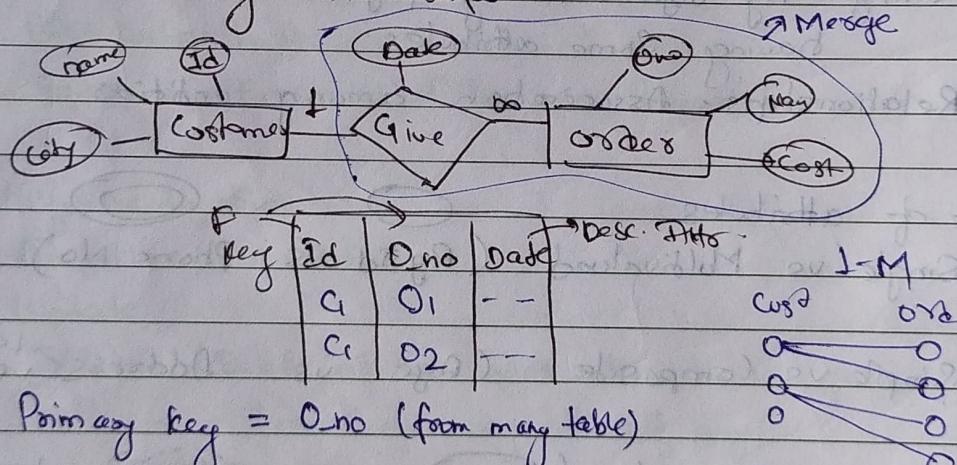
= Prey is any of the foreign key in 1-1 relationship

= We can merge the non-primary attribute column to the table from where Prey was considered.

e.g. ~~Eid Ename age Dept~~ }
 Prey ~~Eid~~
 (in relation)
 ↓
 P1 } Reduction

* Lower the cardinality, higher the duplicacy (1-M) ^{deletes}

(ii) One to Many relationships



2. Reduction of # of tables

Merge orders with Give

Id O no item_name cost Date

9 O1

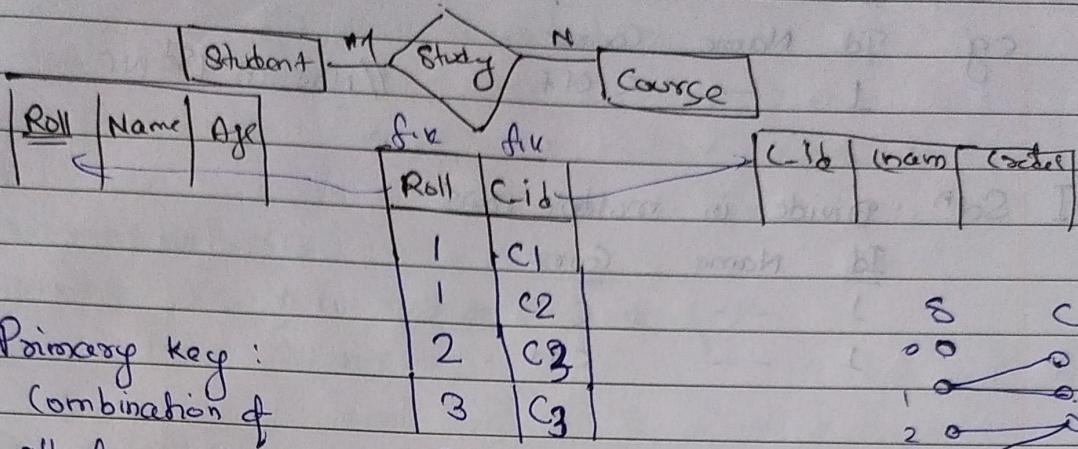
9 O2

C2 O3

C3 O4

Pkey

iii Many to Many



Primary key :
Combination of

all f. keys is the P-key (composite key)

Non-reducible.

Normalization

A technique to remove or reduce redundancy from a table.
→ Row level → Column level.

Insertion, Deletion and update anomaly.

STD	Sname	Grd	Cname	Ad	Fname	Salary
1.	--	C ₁	Math	F	--	304
2.	--	C ₁	Math	F	--	304

(i) Inserting a new course in this table will create a problem as Pkey is SID and initially, no student has taken the newly introduced course.

(ii) Deleting a student will also delete, course details.

(iii) Updation of a soft. colm val (like Salary) will consume extra amnt of time, as it will have to reflect changes in all the similar values.

Soln: Divide the tables:

First Normal Form

Table should not contain multivalued attributes

e.g. Id Name Course

1 -- C, C++

// Not in 1st NF.

I Sdn : Divide in multiple rows

<u>Id</u>	Name	<u>Course</u>
1	--	C
1	--	C++

II Id Name Course1 (Course2)

III Base Table Referencing Table

<u>Id</u>	Name	<u>Id</u>	Course
1	--	1	C
2	--	1	C++

Closure Method (To find all possible candidate keys)

R(ABCD)

FD $\Delta A \rightarrow B, B \rightarrow C, C \rightarrow D$

$A^+ = ABCD$

$B^+ = BCD$

$C^+ = CD$

$D^+ = D$

$C^+ = \{A\}$

Prime attribute = {A}

Non-prime attribute = {B, C, D}

e.g.

R(ABCD)

FD = $\Delta A \rightarrow B, B \rightarrow D, E \rightarrow C, D \rightarrow F$

Right side $\{BDC\}$

$\hookrightarrow F$ not there, \Rightarrow for this

C-key E must be there in LHS -

$E^+ = EC$ X

$AE^+ = BECDFA$ ✓

C-key = $\{AE\}$

check if A is there in RHS

$BE^+ = ABDE$ ✓

$CE^+ = BCD$ X

$D^+ \Rightarrow C\text{-key } \{AE, DE\} \Rightarrow C\text{-key } \{AB, DE, BC\}$

Functional Dependency:

→ Trivial F.D

If $X \rightarrow Y \Rightarrow Y$ is a subset of X .and $L.H.S \cap R.H.S \neq \emptyset$ e.g. $AB \rightarrow A \Rightarrow A = A$ or $A \subseteq AB$

→ Non-trivial F.D

If $X \rightarrow Y$, Y is not a subset of X . $L.H.S \cap R.H.S = \emptyset$

Properties of FD

(i) Reflexivity: if Y is a subset of X , then $X \rightarrow Y$ (ii) Augmented: if $X \rightarrow Y$, then $XZ \rightarrow YZ$ (iii) Transitive: if $X \rightarrow Y$ and $Y \rightarrow Z \Rightarrow X \rightarrow Z$ (iv) Union: if $X \rightarrow Y$ and $X \rightarrow Z \Rightarrow X \rightarrow YZ$ (v) Decomposition: if $X \rightarrow YZ \Rightarrow X \rightarrow Y \& X \rightarrow Z$ (vi) Pseudo-transitivity: if $X \rightarrow Y$ and $WY \rightarrow Z$ then $WX \rightarrow Z$
if $X \rightarrow Y$ and $Z \rightarrow W$ then $(XZ \rightarrow YW)$

Validity is checked for non-trivial FDs:

Second Normal form

1. Table must be in First Normal Form

2. Non partial dependency is all the non-prime attributes should be fully functional dependent on candidate key.

i.e. If composite p.key is there, non-prime attributes should be fully dependent on them, and not any part of the composite key
(partial dependency)

#

Soln. Divide the table.

Steps to check whether a given relation is
2nd NF.

Find C-key, Prime Attribute, Non-prime attributes.

Check for each NPA, if it is determined by

a proper subset of Ckey, then it's not in

2nd NF. OR LHS is a proper subset of Ckeyand RHS is a NPA (then Not in 2nd NF)

Third NF

→ Table must be in 2nd Normal Form.

→ No transitive dependency should be there.

(i.e. A NPA should not determine a NPA).

For each P.D. Check if L.H.S must be
a Ckey or Skey OR R.H.S must be Prime attr.

Find Candidate key:

 $R(ABCD)$ P.D. $A \rightarrow BC$, $D \rightarrow A$ (i) Check R.H.S $\{C, D, A\}$

B is not there, we need B in L.H.S

to form a key.

$$B^+ = \{B\}$$

$$(AB)^+ = \overbrace{AB}^{Reflexive} CD$$

$$\text{Ckey} = \{ABC\}$$

if found (ckey):

check if A or B appears in R.H.S

if yes replace with their determinant

$$AB \rightarrow (DB)^+ = \cancel{DBAC}$$

 DB $\hookrightarrow AB$ (already done)

BCNF (Boyce Codd Normal Form)

→ Table must be in 3NF

→ For each FD: L.H.S must be a key or Superkey



DBMS

continued after normalization:

Lossless and lossy Decomposition

Whenever we tend to normalize the relation, we decompose them into one or more relations.

In the process of again joining them, if there is a loss (inconsistency) of data then it was lossy decomposition else lossless.

→ During decomposition, at least one common attribute should be there. And the common attribute should be very/very few either R_1/R_2 or both (unique).

After joining the decomposed relation; even if there are extra tuples it is lossy decomposition.

The extra tuples which appears in the joined table, they are called Spurious table tuple.

e.g. R			R ₁		R ₂	
A	B	C	A	B	A	C
1	2	1	1	2	1	1
2	2	2	2	2	2	2
3	3	2	3	3	3	2

R₁ Natural join R₂ gives

A	B	C	(remove tuples where $R_1.A \neq R_2.A$)
1	2	1	X
1	2	2	X
1	2	3	X

2 2 1 1 X

2 2 2 2

2 2 3 2 X

3 3 1 1 X

3 3 2 2 X

3 3 3 2

A B C

1 2 1

2 2 2

2 2 3

3 3 2

Lossless decomposition

$\therefore A$ was very few

$R_1 \& R_2$ both

(No. Spurious tuple)

Conditions of lossless decomposition:

$$(1) R_1 \cup R_2 = R$$

$$(2) R_1 \cap R_2 \neq \emptyset$$

(3) The common attr. must be c.keys/ key of either of decomposed tables or all tables.

Fourth Normal Form

\rightarrow BCNF + No multivalued dependency

4th Normal Form

Name	Mobile	Email
Vasun	M1	E1
Vasun	M1	E2

→ Multivalued dependency

" M2 E1

" M2 E2

Vasun M3 E1

" M3 E2

$x \rightarrow y$

~~mult~~

\rightarrow Name \rightarrow Mobile

Name \rightarrow Email

5th Normal Form

4th N.F + Lossless Decomposition

Minimal cover in DBMS

Step 1: In R.H.S only one key cells

Step 2: Remove redundant

Step 3: AC \rightarrow B

\rightarrow C+ has then C+D \subseteq AC+D

L.H.S to one attr.

e.g. find minimal cover of $\{A \rightarrow B, C \rightarrow B, D \rightarrow ABC, AC \rightarrow D\}$

Step 1: $A \rightarrow B, C \rightarrow B, D \rightarrow A, (D \rightarrow B), D^+, A^+, D$

Step 2: check if $A \rightarrow B$ can be removed $A^+ = A$ X

" " $C \rightarrow B$, $C^+ = C$ X

$D \rightarrow A$, $D^+ = DBC$ X

$D \rightarrow B$, $D^+ = DABC$ ✓

$D \rightarrow C$, $D^+ = DAB$ X

$AC \rightarrow D$, $(AC)^+ = ACB$ X

We get $A \rightarrow B, C \rightarrow B, D \rightarrow A, D \rightarrow C, AC \rightarrow D$

Step 3:

$$\begin{array}{l} AC \rightarrow D \\ \times \\ CT = CB \quad (\text{No } A) \end{array}$$

$$AC \rightarrow D$$

$$\begin{array}{l} \times \\ AT = AB \quad (\text{No, C}) \end{array}$$

Final ans;

$$\begin{array}{l} \downarrow \text{Union} \\ AC \rightarrow B, C \rightarrow B, D \rightarrow AC, AC \rightarrow D \end{array}$$

Q. Check highest NF for a given func Dep.

→ Find Candi. by

→ Find Prime & NPA

→ For each FD check from BCNF to 1st NF.

If all satisfies then

(Table is assumed to be in 1st NF)

II Finding out NF of a relation

$$\text{e.g. } R(ABCD E F P) \quad CR = \{AB, FB, EB, (B)\}$$

$$PA = \{A, B, C, E, F\} \quad NPA = \{D\}$$

$$FDs \{AB \rightarrow C, C \rightarrow D, (\rightarrow E, E \rightarrow F, F \rightarrow A)\}$$

III Remove partial dependency. { L.H.S must be proper subset
of R.H.S
Candi. of C.R. }
 \hookrightarrow
 $C \rightarrow D$ (NPA)

proper subset

R.H.S must be a NPA.

Now divide the table & (in a lossless way)

$$(1) R(ABCDEF)$$

$$\begin{array}{l} AB \rightarrow C \\ C \rightarrow E \\ E \rightarrow F \\ F \rightarrow A \end{array}$$

$$CR \{AB, PB, EB, (B)\}$$

$$R_1(ABCDEF) \quad \text{and} \quad R_2(CD)$$

$$\begin{array}{l} CT = CD \\ (No \text{ no PD}) \end{array}$$

Aug & C.R.

S-II Remove transitive dependency

\rightarrow L.H.S is crg / strg & R.H.S is a PA (Not in TD)

(R₁) $AB \rightarrow C, C \rightarrow E, E \rightarrow F, F \rightarrow A$

(R₂) $C \rightarrow D$

Already in 3rd NF.

S-III check for BCNF

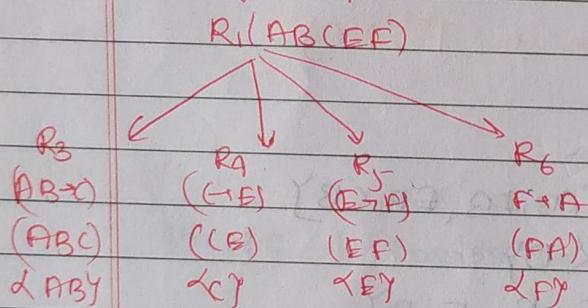
\rightarrow L.H.S must be crg, R.H.S must be PA

R₁ $\nsubseteq AB \rightarrow C, C \rightarrow F, E \rightarrow F, F \rightarrow A$

Not in BCNF

R₂:
 $C \rightarrow D$

In BCNF
(0% redundant)

Decompose R₁

R₂ (\rightarrow D)

R₂ (D)

dcy

Lossless join:

Let's say a query about B & F

(combine R₃ & R₇ (c is crg))

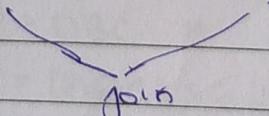
But if we want to join R₃ (ABC) & R₇ (EF)

\rightarrow we can't join directly;

first join R₃ & R₈ then with R₇ (EF)

(ABC₋E)

(EF₋)



Date: 08/08/21

Equivalence of functional dependency

Given $X \& Y$ check X covers Y $X \geq Y$ Y covers X $Y \geq X$ $\Rightarrow X \equiv Y$ (equivalent)(i) check if X covers Y eg $X = \{A \rightarrow B, B \rightarrow C\}$, $Y = \{A \rightarrow B, B \rightarrow C, A \rightarrow C\}$ for each FD in Y , take closure from X to see if the closure covers the determination from Y .(i) for $A \rightarrow B$ (from Y)

$$A^+ = ABC \quad (\text{closure from } X)$$

$\hookrightarrow B$ covered

(ii) $B \rightarrow C$

$$B^+ = BC \quad \begin{matrix} C \\ \hookrightarrow \text{covered} \end{matrix}$$

(iii) $A \rightarrow C$

$$A^+ = ABC \quad \begin{matrix} C \\ \hookrightarrow \text{covered} \end{matrix}$$

So, T & T & T $\Rightarrow X$ covers Y .

(i) (ii) (iii)

Sly, check whether Y covers X (i) $A \rightarrow B$ (X)

$$A^+ = ABC \quad \begin{matrix} C \\ \hookrightarrow \text{covered} \end{matrix}$$

(ii) $B \rightarrow C$

$$B^+ = BC \quad \begin{matrix} C \\ \hookrightarrow \text{covered} \end{matrix}$$

So, Y covers X $\Rightarrow X \equiv Y$

Dependency preserving decomposition

$\Rightarrow R$ is decomposed into R_1 & R_2
 $\Rightarrow FD_1 \cup FD_2 = FD^+$ (then dependencies are preserved)

e.g. $R(ABCD)$

$$FD = \{A \rightarrow B, B \rightarrow C, C \rightarrow D, D \rightarrow B\}$$

decomposed into $R_1(AB)$, $R_2(BC)$, $R_3(CD)$

$R_1(AB)$	$R_2(BC)$	$R_3(CD)$
$A \rightarrow A$ ✓ found (exists in R)	$B \rightarrow C$ ✓	$B \rightarrow D$ ✓
$A \rightarrow B$ ✓	$C \rightarrow B$ ✗	$D \rightarrow B$ ✗
$B \rightarrow A$ ✗ doesn't exists	$C \rightarrow CD$ ✗	$D \rightarrow CD$ ✗
$B \rightarrow A$ (not found)		
$\rightarrow FD^+ = BCD$ ((loss from R))		

So final union of FDs
 $FD^+ = FD_1 \cup FD_2 \cup FD_3 = \{A \rightarrow B, B \rightarrow C, C \rightarrow B, B \rightarrow D, D \rightarrow B\}$

(closure in FD^+) $CF = CBD$ ($C \rightarrow D$)

So, dependencies are preserved.

+ Joins & Subqueries gives similar results, later is complex.

Joins = cross product + Select statement (condition)

\rightarrow Cross join (cross product)

\rightarrow Natural join

\rightarrow Conditional

\rightarrow Equi join

\rightarrow Self join

\rightarrow Outer

full
left right

Natural join

e.g. find emp name who is working in a department

emp

eno	ename	addr
-----	-------	------

dno	Name	cn
-----	------	----

Ans: Select ename from $(\text{Emp} \bowtie \text{Dept})$, cross product
where $\text{Emp.eno} = \text{Dept.eno}$;



natural join

Select ename from $(\text{Emp} \text{ Natural join } \text{Dept})$

In Natural join

cross product is calculated; and only the tuples where values for common attribute are same are taken into consideration rest are ignored.

e.g. In above example, eno was common attribute.

Self-Join : Table joined with itself

(SQL query starts with from)

Study

sid	cid	since
s1	c1	2016
s2	c2	2017
s1	c2	2017

e.g. find student id who is enrolled in at least two courses

Select T1.sid from Study as T1, Study as T2

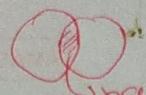
where $T1.sid = T2.sid$ and $T1.cid <> T2.cid$;

In Self-Join

not equal to

⇒ Cross product with itself

⇒ Apply some equivalent (non-equivalent constraints according to need)



natural join

Equi-join

e.g. Find the employee whose department's location is same as their address

Emp			Dept		
eno	ename	address	dno	location	eno

SQL query : Select ename from emp,dept
where emp.eno = dept.eno and
emp.address = dept.location;

In equi-join :

Natural join + some additional constraint (equivalence)

Left outer join :

It gives the matching rows and the rows which are in left table but not in right table.

Emp			Dept		
eno	ename	dno	dno	ename	location

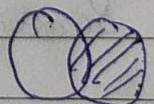
eg. Select eno,ename,dname,loc
from emp left outer join dept
on (emp.dno=dept.dno);

$$\text{for } A, B \Rightarrow A \cap B + (A - B)$$

Right outer join :

It gives the matching rows and the rows which are in right table but not in left table.

$$\text{for } (A, B) = A \cap B + (B - A)$$



In general, left / right / full
= Natural join + something

... emp right outer join dept ...

Full outer join = (Left outer join) \cup (Right outer join)

Relational Algebra (formal query language) (procedural language)

↙
what to do? + how to do?

Operations

Basic operation

- Projection (Π)
- Selection (σ)
- Cross product (\times)
- Union (\cup)
- Rename (ρ)
- Set Difference ($-$)

Derived operation

- Join (\bowtie)
- intersect (\cap)
- $(X \cap Y) = X - (X - Y)$
- Division ($/, \div$)

Projection (Π)

Query: Refine the roll number from Student (table)

$\Pi_{\text{Rollno.}}$ (Student)

→ It projects only distinct values (no duplicate data in a column is projected)

→ Projection is used at end.

Selection (σ)

Query: Refine name of student, whose Rollno=2.

$\Pi_{\text{name}}(\sigma_{\text{Roll-no}=2}(\text{Student}))$ (Row level selection)

→ narrows search area.

↓
selects that particular row
(projection at end)

column level projection

Date: 01/08/21

(Student - product) (At least 1 col should be common)
In answer

No. of columns = Sum of columns

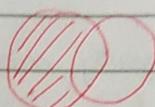
No. of Rows = Product of rows

$R_1 \times R_2$

Set difference (-)

$A - B = A$ but not B

= $A - (A \cap B)$



1) No. of columns must be same

2) Domain of every column must be same.

eg	Roll	Name	Eno	Ename
	1	A	7	E
	2	B	1	A
	3	C		

(Student - Employee) = Roll name

2	B
3	C

$T_{name}(\text{Student}) - T_{name}(\text{Employee})$

1	name
	B
	C

Union

$A \cup B = A + B - A \cap B$

→ conditions same as set difference

Roll No. Name -

$(\text{Student} \cup \text{Employee}) =$	1	A
	2	B
	3	C
	7	E

$T_{name}(\text{Student}) \cup T_{name}(\text{Employee})$

Division operation

for 'ALL' and 'EVERY', use SQL division

$$A(x, y) / B(y) =$$

Find sid who are enrolled in all subjects $\rightarrow \text{ops}(s)$

from
whose
class
is
expected

Enrolled		Course
Sid	Cid	Cid
s1	c1	c1
s2	c1	c2
s1	c2	
s3	c2	

→ find cross product of sid (enrolled) \times course

→ Set difference of above s/p and enrolled

→ in cross product but not in enrolled

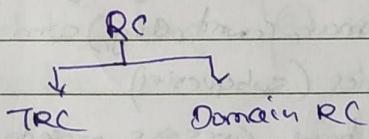
↳ this gives disqualified candidates

→ Total - disqualified = qualified (studying all subjects)

$$T_{sid}(\text{enrolled}) - \left(T_{sid}(\text{enrolled}) \times T_{cid}(\text{course}) - T_{sid}(\text{enrolled}) \right)$$

$$\text{Total} - \text{disqualified}$$

Tuple Relational calculus (Non-procedural)



Query is expressed as

$$\alpha t \mid P(t)$$

↳ tuples satisfying predicate $P(t)$ [conditions]Operations: OR(\vee) , AND(\wedge) , NOT(\neg)Quantifiers: $\exists t \in \gamma (Q(t))$, if any true \Rightarrow true $\forall t \in \gamma (Q(t))$, is true for all tuples in relation γ

SEQUEL: simple eng. query lang.

AMAS Page No.

EE code: father of DBMS.

Date

Unsafe expression:

e.g. $\exists Sname (\neg \text{Supplier}(S))$

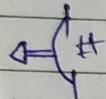
↳ gives ∞ .

Free variable:

\exists, \forall is not used

Bounded

\exists, \forall is used before



SQL

IBM implemented the proper published by EE code.

90% of data is non-structured, so NOSQL is
in trend.

→ SQL is domain-specific language

(we can only use it in structured / relational
data)

→ SQL is a declarative language

(only what to do)

But PLSQL is procedural as well.

Operators (like B/w, In, Not in, Conditional)

Clauses distinct, order by, group by, from, having

Aggregate func (min, max, count, Avg) sum

Joins & Nested Queries (subqueries)

PLSQL (Triggers, functions, cursors, procedures)

SQL commands

DCL: Create, Alter, Drop, truncate, Rename

DML: Insert, update, delete, select. (CRUD)

DCL : Grant, Revoke

TCL : Commit, Rollback, Save point

Constraints: Primary, Foreign, Check, Unique, Default, Not Null.

Create Table emp (

 column1 type, . . .) ;

desc tablename; // Shows structure of table

Number (scale, precision)

ALTER Command:

→ Add column

 Alter table name add (address varchar(30));

→ Remove (column(s))

 Alter table name ~~drop~~ ~~remove~~ column colname;

→ Modify datatype

 Modify colname datatype (new)

→ Modify datatype length

 Add Primary key (rollno);

→ Add / Remove constraints

 Rename column colnam ^{to} newnam;

→ Rename column / Table.

 Rename to newtablename;

Difference b/w Alter and update

(1) DDL

(6) DML

(2) Used to change
structure

(2) Used to change tuple's
data.

e.g. update tablename set

 salary = salary * 2;

Diff b/w Delete

(1) DML

Drop

Truncate

(2) Deletes tuples

DDL

DDL

eg Delete from tablename
where id = 2;

Deletes

Truncate student;

(It is a special
case of delete)

(3) Slower

else

(3) Faster

(4) Rollback

else tab name

(4) Rollback not possible

possible

doesn't exist

(if committed)

(2) No rollback

Constraints in SQL

(Restrictions before putting the data)

- 1) Unique
- 2) NOT NULL
- 3) Primary Key (Unique + Not Null)
- 4) Check
check (age > 18)
- 5) Foreign Key
- 6) Default
e.g. Salary int default 1000;

SQL queries and Subqueries

Emp

eid	ename	Dept	Salary
1	Ram	HR	10000
2	Amit	MRKT	20000
3	Ravi	HR	30000
4	Nitin	MRKT	40000
5	Vasun	IT	50000

Q1. Display employ name who is taking max salary

Select ename from Emp

where salary = Select max(salary) from Emp;

Inner query

Date: 05/08/21

Q2 Find 2nd highest salary & taking employee's name.

→ Inner queries runs for only one time

Select ename from emp where salary

(Select max(salary) from emp where

salary <> (select max(salary) from emp));

Q.3: Write a query to display all the dept names along with no. of emps having working in that.

Group by to group similar things together.

(i) Only the column after 'group by' can be written after 'select'.

(ii) Apart from above aggregate func. can be used
Select dept, count(*) from emp group by dept;

MRKT	2
HR	2
IT	1

Q.4: Write a query to display all the names of dept where no of emps are less than 2.

Instead of 'where' we use having with group by

Select dept from emp group by dept
having count(*) < 2;

Q.5 Display highest salary dept, and name of emp who is taking that salary

Select e.name, from emp where

salary in (Select max(salary) from emp

// group by dept);

equal to won't be used (\because this query has multiple of rs)

Project

eid	pid	Pname	Location

IN / NOT IN

from
Select * , emp where

address IN ('Chandigarh', 'Delhi');

Find the name of the employees who are working on a project

Select name from emp whose eid IN (Select eid from project);

Exists/Not exists

Correlated query (Top down)

For each outer query iteration, run the inner query

g. Find the details of emp who is working on at least one project.

Select * from emp where eid exists (Select eid from project where emp.eid = project.eid);

It behaves like two nested for loops in C language.

SQL aggregate functions

Count(column) won't count a tuple if a null value is there.

Sum

DISTINCT(Sum(salary))

Sums counts distinct value only.

Avg.

$\text{Ave}(\text{salary}) = \frac{\text{Sum}(\text{salary})}{\text{Count}(\text{salary})}$

Distinct(Ave(salary)) = Distinct(Sum(salary))

Distinct(Count(salary))

Correlated Subquery (Synchronized query)

↳ It is a subquery that uses value from outer query (top-down approach)

(Q.) Find all employees who work in a department

Emp	Dept
eid ename add.	Did Dname eid

Select * from emp where

Dept exists (Select * from dept

where emp.eid=dept.id);



Correlate outer query value \Rightarrow entire row

Nested Subquery

Bottom-up

Correlated query

Top-down

Joins

(Crossprod + constraint
attribute)

Select * from emp
where eid in
(Select eid from
dept);

Select * from emp
where eid exists
(Select id from dept
where emp.id=dept.id);

Select * from
emp, dept
where emp.id = dept.id;

\Rightarrow less no. of
comparisons

exists returns

True or False

\rightarrow works faster
but takes space

Find Nth highest salary using SQL

Select id, salary from emp e1

where N-1 = (Select count(Distinct salary) from

emp e2 where

e2.salary > e1.salary);

No. of salary greater than e1.sal in e2.

For 4th highest, 3 higher salaries should be
there in e2.

LIKE command

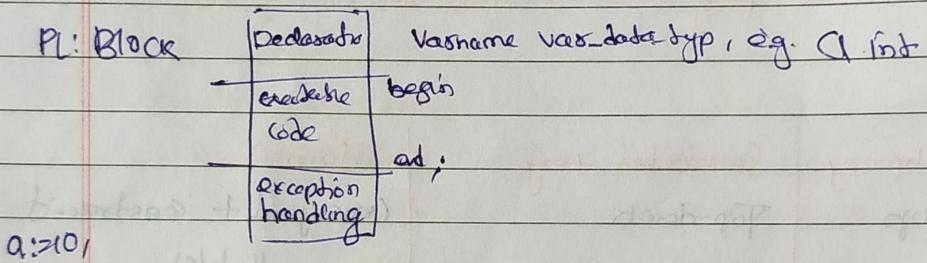
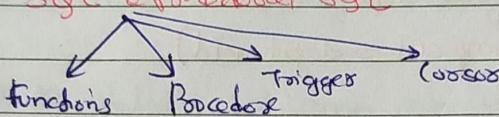
- % → anything
- _ → places = no. of underscores

→

(Select * from R) intersect (Select * from R)
→ (produces unique tuples from *)

PL-SQL & Procedural SQL

(what & how to do both)



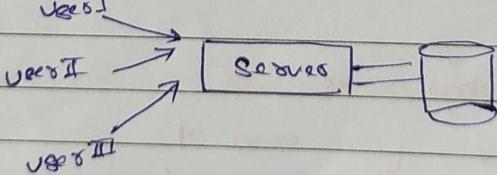
Transaction Concurrency

Transaction : A set of operations used to

perform a logical unit of work.

→ It generally represent change in database.

⇒ Basic operations in DB → Read/Write or commit/rollback



ACID properties

Atomicity : Either all or none

Consistency : Before the transaction starts and
after the completion of transaction

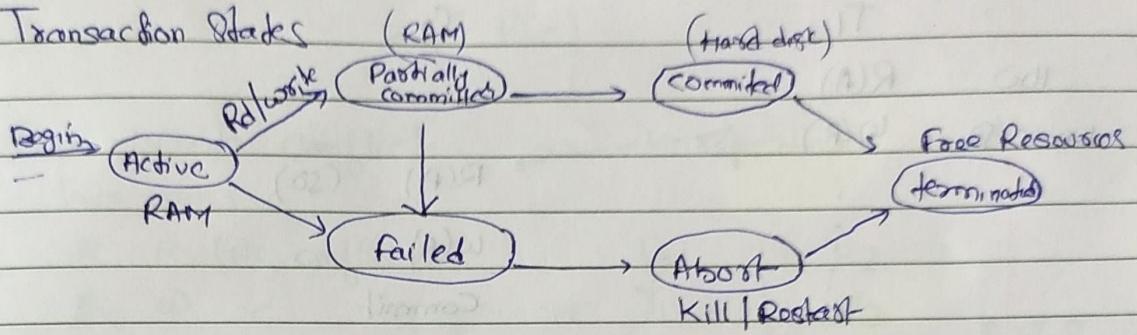
Sum of money should be same.

Isolation:

Converting a parallel schedule into serial schedule.

Durability: All changes made to DB should be permanent.

Transaction States



If 'n' instructions are there

- (n-1) use general instructions $\xrightarrow{\text{complete}}$ partially committed
- nth instruction is 'commit' \rightarrow committed.

Failed/Aborted transactions can't be resumed, it has to be started.

Schedule: chronological execution sequence of multiple transactions

Serial,

$T_1 \rightarrow T_2 \rightarrow T_3$

T_1

T_2

T_3

once a transaction is not done with execution, no other trans. can start

→ consistency

→ Large waiting time

→ Low throughput

Parallel

T_1

T_2

T_3

T_2

T_1

→ Good performance

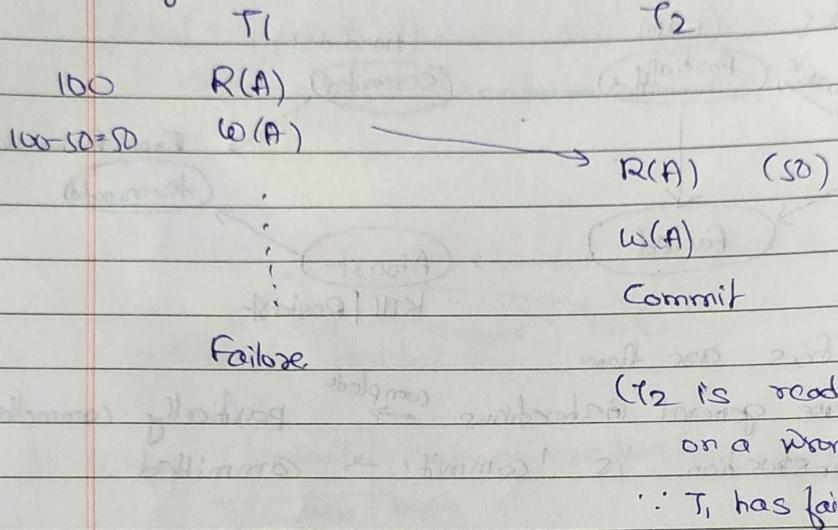
→ High throughput

→ Better resource utilization

Types of problems in Concurrency

- Dirty read → Lost update
- Incorrect Summary → Unrepeatable read
- Phantom read.

Dirty read



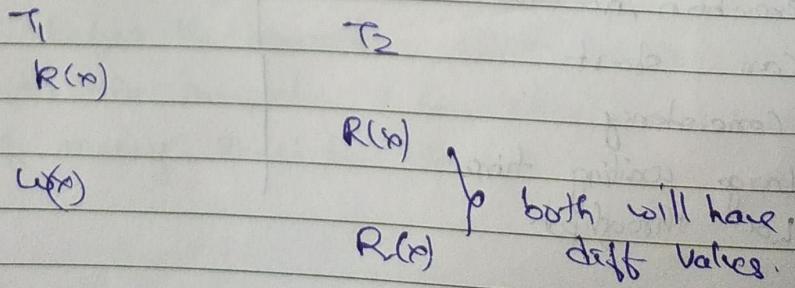
Incorrect Summary:

While one transaction is performing some operations, and meanwhile another transaction T_2 enters and starts executing, and performing some agg. functions; this will always lead to wrong results.

Lost update

If two transaction tries to update the same data; by reading simultaneously, the data written by first committing transaction will be lost.

Unrepeatable read



Phantom Read

 T_1 T_2 $R(x)$ $R(x)$ Delete (x) $R(x)$

Read - write Conflict or Unrepeatable read

Four cases

 $R \ R$ (No conflict) \times $R \ W$ ✓ $W \ R$ ✓ $W \ W$ ✓ U_1 U_2 T_1 T_2 $A \neq 0$ 2 $R(A)$ $R(A) = 2$ $w(A) \neq A = A - 2$

(commit)

No of books
↓
 $A = 0 \neq 9$ T_1
10 $R(A)$
 $A = A - 1$ T_2 $R(W)$ problem0 $R(A)$

9

 $R(A) 10$ A = A - 1 9 $w(A)$

(commit)

 $w(A)$

commit

⇒ 2 books issued but only one decrement $\Rightarrow R/W$ problem

Date: 06/08/21

 $A = 0 \neq 10$

Recoverable Schedule

Here even if

→ has completed

it's transaction

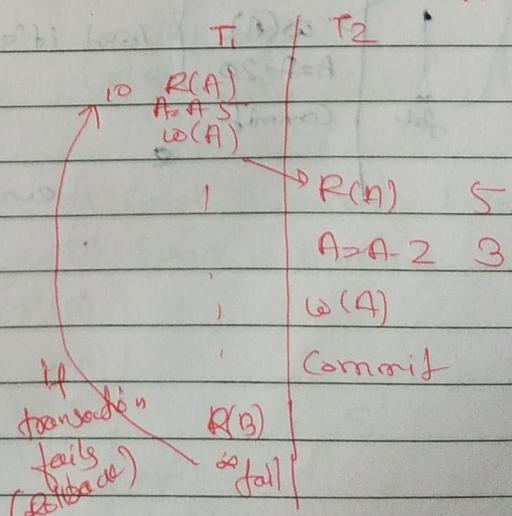
Database roll backed

so it's initial state.

we lost all changes

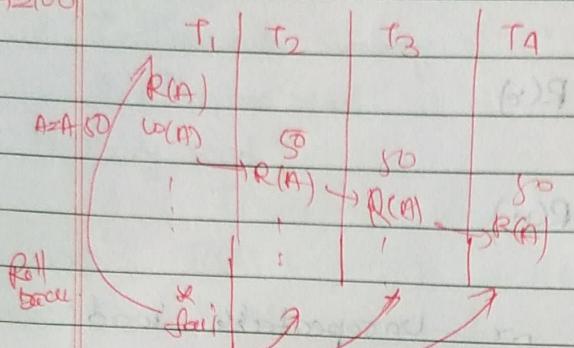
due to T_2 .

It is irrecoverable.



Cascading vs Cascadless Schedule

A=100



An rollback if one fails

Disadvantage:

Performance of CPU degraded

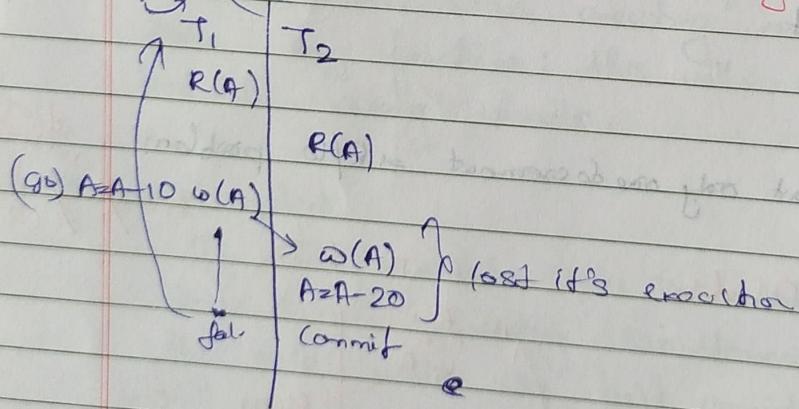
Waste of time

We can make it cascadelss

- by stopping other transactions from reading till it is not committed to db.
i.e. T₂, T₃, T₄ can't read (lock); till T₁ has been committed.

Work-work problem in Cascading Schedule

A=100(RB)



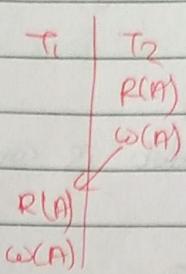
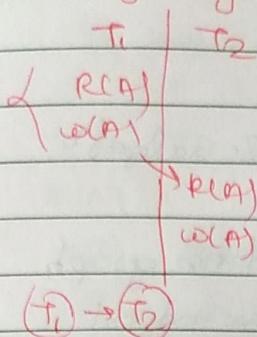
even if T₂ has committed

80; T₁ will

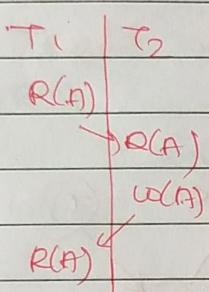
roll back and

re-initialize db values

Serializability



Both case serial, no need to
Serialize



We need to find a clone for
this schedule which is serial.
Then it's serializable.

Parallel

Finding Conflict equivalent

$R(A)$ $R(A)$ } Non-conflict pairs

$R(A)$ $w(A)$

$w(A)$ $R(A)$

$w(B)$ $w(A)$

$R(B)$ $R(A)$

$w(B)$ $R(A)$

$R(B)$ $w(A)$

$w(A)$ $w(B)$

S

T_1 | T_2

$R(A)$ |

$w(A)$ |

S'

T_1 | T_2

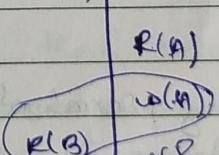
$R(A)$ |

$w(A)$ |

$R(B)$ |

$R(B)$

$w(A)$



T_1 | T_2

$S \Rightarrow R(A)$

$w(A)$

$R(A)$

$R(B)$

$w(A)$

(i) Find non-conflict pairs

(ii) Swap in timeline

$R(A)$

$w(A)$

$R(B)$

$= S'$

$R(A)$

$w(A)$

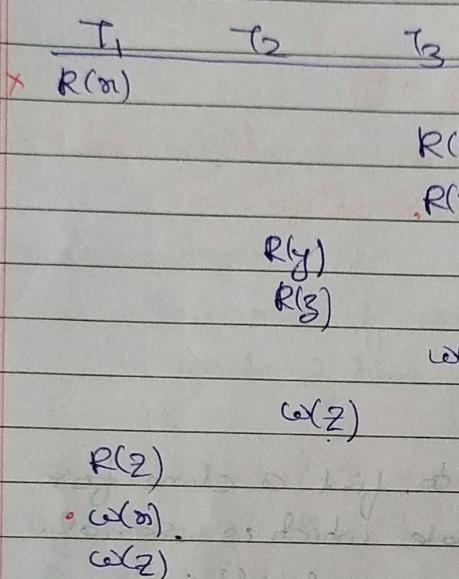
$\Rightarrow S \leq S'$

(Conflict equivalent)

$S \xrightarrow{E} S' \rightarrow \text{Serializable}$

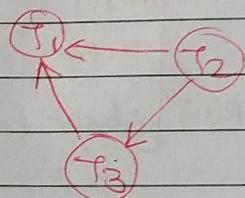
Conflict Serializability

- check if a schedule is conflict serializable



- precedence graph
(To draw edges)

- choose conflict pairs in other transactions and draw edges



e.g. for $(R(a))$ we find

~~$R(b) \omega(a) m$~~

remaining transaction..

- check if a loop exists

if NO loop

→ this schedule

is conflict serializable.

we can find conflicts

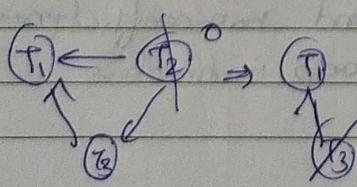
on serialized schedule

→ It is consistent

finding serialized schedule out of all possibilities

(i) find indegree of all nodes, start with

zero indegree node remove and add to answer



$T_2 \rightarrow T_3 \rightarrow T_1$

No loop exists

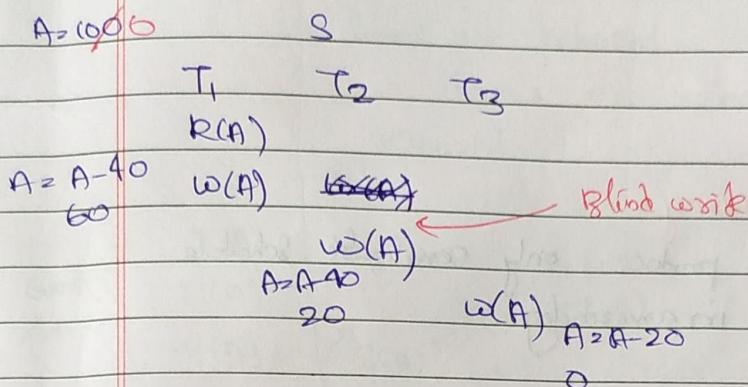
conflict serializable

Serial

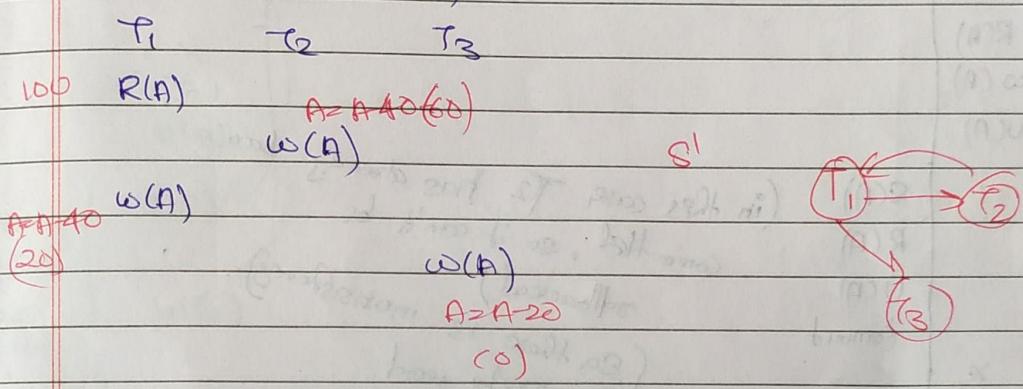
consistent

View Serializable

If loop is there, we need to check for view serializable.



Another scheduling (S')



Here, both S and S' give same result although.

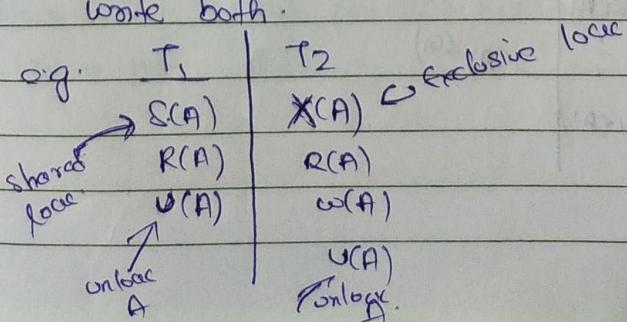
S' has a loop and it's non-conflict serializable
but \neq S which is serial and equivalent to S'
yielding same results.

1. # Shared-exclusive Locking

T1 shows how to make them recoverable/serialized.
They are known as concurrency control protocols.

Shared Lock (S) \Rightarrow if transaction locked data item in shared mode then allowed to read only.

Exclusive Lock (X) \Rightarrow if transaction locked data item in exclusive mode then allowed to Read and write both.



Compatibility table:

	S	X
S	✓	✗
X	✗	✗

Problems in S/X locking:

- (i) May not be sufficient to produce only serializable schedule
- (ii) May not be free from irrecoverability
- (iii) T₁ | T₂

	T ₁	T ₂
Rollback	X(A)	
	R(A)	
	W(A)	
	U(A)	
!	S(A)	(in this case T ₂ has already been committed, so it can't be rolled back)
,	R(A)	
,	U(A)	(so, there is inconsistency)
!	commit	
x		only read
fail		

- (iv) May not be free from deadlock.

	T ₁	T ₂
(G)	X(A)	
(wait)	X(B)	(G)
		X(A) (wait)

- (iv) May not be free from starvation
- a transaction has to wait due to a more prior transaction, but not infinitely.

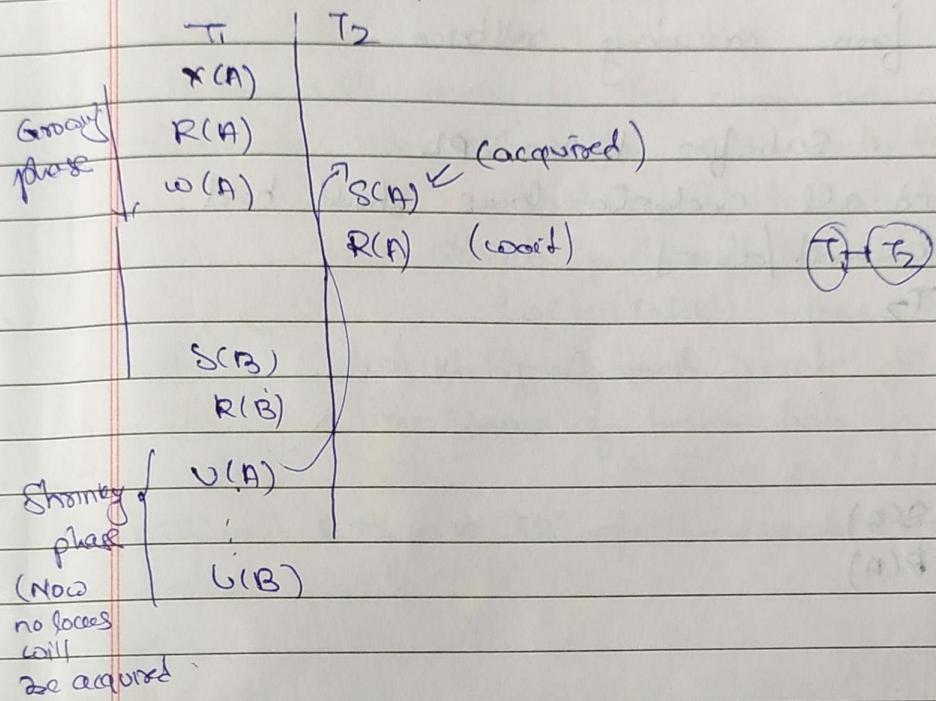
	T ₁	T ₂	T ₃	T ₄
		S(A)(G)		
(G)	X(B),		(G) S(A)	
		U(A)		S(A)(G)

T₁ is starving of

2-phase Locking (2PL)

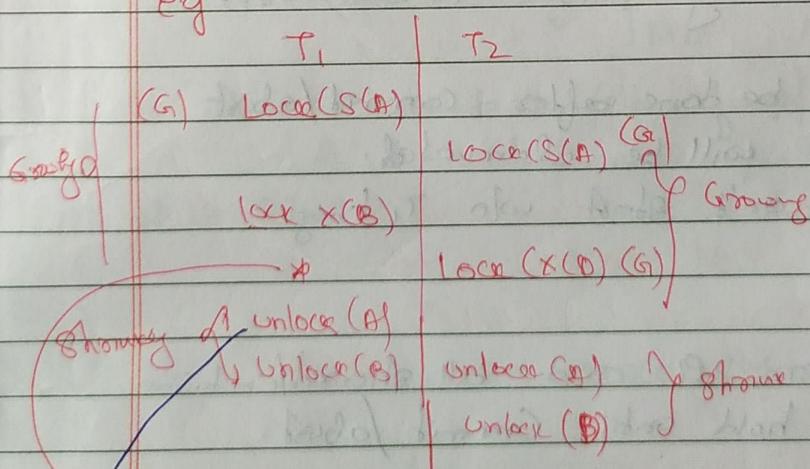
Growing Phase: locks are acquired and no locks are released.

Shrinking Phase: locks are released and no locks are acquired.



T_2 's growing phase may also start while T_1 's growing phase is going on based on compatibility.

e.g.



Lock Point is last pt. where T_1 was locked

Unlock pt.: first " " " " " unlocked.

2PL always ensures serializability

Drawbacks in 2PL

- (1) May not free from nonrecoverability
- (2) Not free from deadlocks
- (3) Not free from starvation
- (4) Not free from cascading rollback

Strict 2PL & Satisfies basic 2PL

and all exclusive locks should hold until commit/abort.

T_1	T_2
$X(A)$	
$R(A)$	
$W(A)$	
	$S(A)$
	$R(A)$
Commit	
$U(A)$	

Note
 $T_2/T_3/T_4$

can read

or ask for $S(A)$, \therefore they will read from db once T_1 commits. So, cascading rollback is eliminated

Unlocking will only be done after commit/abort
 So, inconsistency will be eliminated.

$\therefore T_2/T_3/T_4 \dots$ can't abort w/o T_1 being committed

Rigorous 2PL: Basic 2PL and all shared/exclusive unlockers should hold until Commit/abort.

It has no extra adv. than Strict 2PL.

!! Both the 2PL doesn't talks about deadline and starvation

Conservative SPL: (Practically not possible)

Take all locks on all data before starting.

→ It removes deadlock problem.

Time Stamp Ordering Protocol.

→ Unique value assigned to every transaction.

→ Tells the order (when they enter into the system).

→ Read TS (R_iS): Last (latest) transaction no. which performed 'Read' successfully.

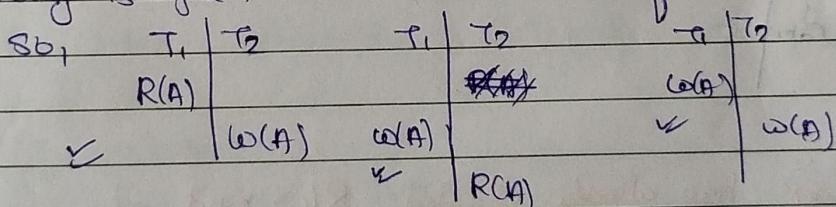
→ Write_TS (W_iT_iS): Last (latest) transaction no. ('3, TS) which performed 'Write' successfully.

→ TS (T_i): Time Stamp of transaction.

WTS and RTS are TS of transaction satisfying criterias

Rules:

Priority is given to older transactions for completion.



∴ we will complete T₁ (older first).

Rule I: Transaction T_i issues a read operation

(a) if WTS(A) > TS(T_i) , Rollback T_i

(b) otherwise execute R(A) operation

~~Rollback~~ set RTS(A) = Max {RIS(A), TS(T_i)}

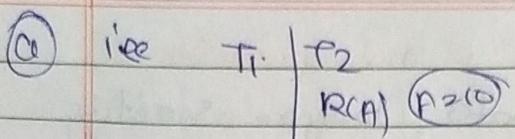
Rule II: Transaction T_i issues W(A) operation

(a) if RTS(A) > TS(T_i), then rollback T_i

(b) if WTS(A) > TS(T_i) , then rollback T_i

(c) Otherwise execute write (A) operation

Set WTS(A) = TS(T_i)

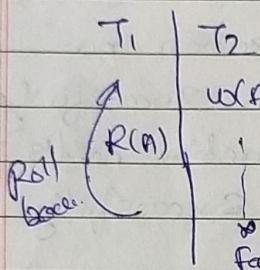


$\text{R} = (20)$ $w(A)$ X won't be allowed
 c

Rule II (a)

Younger has already read but, we will complete first. This won't be allowed

(b) $100 \quad 200$



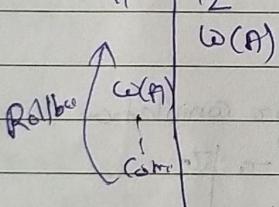
\downarrow
fail

Rollback X

Rule I (a)

don't allow T_1 to read.

(c) $T_1 \quad T_2$



Rule II (b)

If younger has already started R/W and older wants to R/W afterwards, it's not allowed. Since, first older will be completed.

In TSOP $T_1 \rightarrow T_2 \rightarrow T_3$

older —> younger

Indexing

Why indexing is used?

→ It reduces the searching time, because we can search our data according to indices

I/O cost is reduced due to indexing.

			Youngest			
			(200)	(300)		
			T ₁	T ₂	T ₃	
<u>R(A)</u>						A B C
W(C)		R(B)				RIS 0 0 0
R(C)			R(B)			WTS 0 0 0
		W(B) Rollback		W(A)		R(A) for T ₁ WTS(A) = 0 > 100 (for 8) RIS(A) = Max(0, 100) = 100 A B C
						RIS 100 300 0 WTS 300 0 100

Date : 07/08/21

Indexing

With indexing

Eg. Block size = 1000 B

Record size = 250 B \times 10000 (crossed)

As. No. of blocks = 2500 blocks

each block contains $= \frac{1000B}{250B} = 4$ records

$$\text{Avg time} = \frac{\text{worst time}}{2} = \frac{2500}{2} = 1250$$

If ordered

$$\log_2(2500) = 12.$$

With indexing

Eg. Search a record from Index table if
Index table entry = 20B (key + pointer)
 $10B$ $10B$

Block size in Index = Block size in file

Each block in index table contains $= \frac{1000B}{20B} = 50$ Nos.

Dense : For each record in file, we make index

Sparse: Only one index per record (only possible in ordered).

⇒ If sparse, no. of index = 2500 (= no. of blocks)
each block in Index file can hold 50.

$$\text{So, total blocks needed} = \frac{2500}{50} = 50 \text{ Blocks}$$

$$\text{Searching time} = \lceil \log_2(50) \rceil = 6 + \text{H} \quad \text{excess search}$$

(As Sparse ⇒ ordered)

So after 6 search we find second block
in HD then we search those in particular
block for exact record.

(unsorted)

⇒ In dense, no. of indices = 10000

$$\text{total blocks needed in index file} = \frac{10000}{50} = 200 \text{ Blocks}$$

$$\text{Searching time} = \log_2(200)$$

$$= 8 + \text{H}$$

↑ Search record in that
page/block.

Types of indexes

- 1) Primary 2) Clustered 3) Secondary

ordered file	Primary Index (1 2 3 4 5 ...)	Clustered Index (1 1 2 2 3 3 ...)
unordered file	Secondary Index (2 1 3 4 6 7 ...)	Secondary Index (2 2 3 4 ...)

Key Non-key

Primary Index

Conditions

→ Ordered data

→ key (unique)

Try to make R-Index Sparse

each block has one index

ie. # of entries in R-table = # of blocks in HD

(Also, known as Anchor record) → the first entry in each block

1
2
3
4
5
6
7
8

$$\text{Time} = \log_2(N) + 1$$

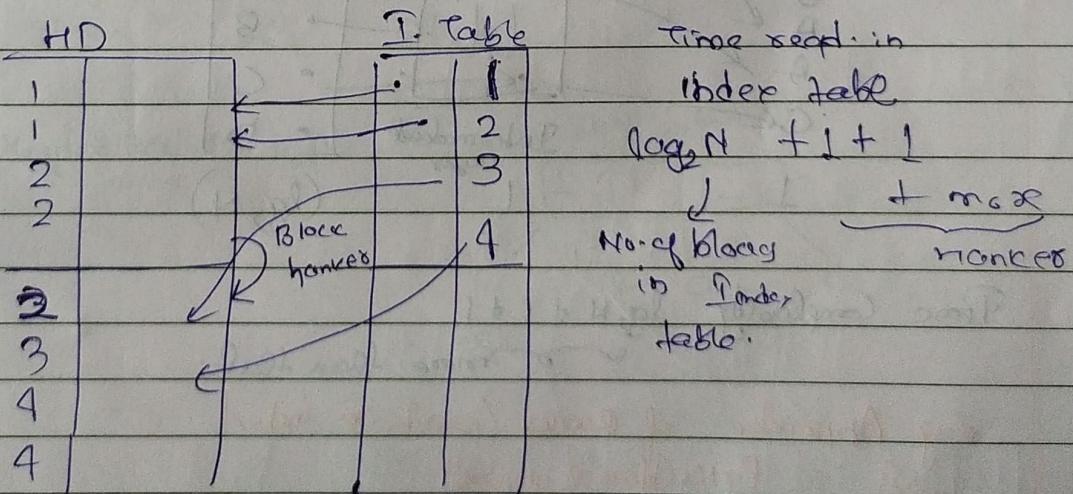
↑
No. of blocks in IT.

This depends on Sparseness.

Clustered Index

→ ordered data

→ Non-key



→ It is also sparse

→

Secondary index (Second index)

→ un-indexed data

→ key / non-key data

Why go Secondary index??

→ if queries are based on some ^{primary} non-key attributes
like Name, phone no.

•

2nd index is always dense. No cases of
un-indexed data.

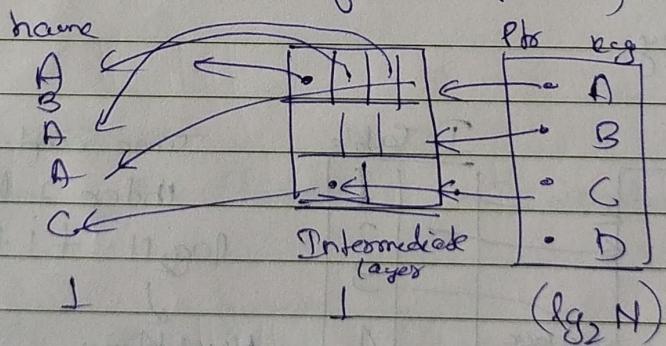
→ All records have individual indices

records (function) = # indices.

If un-indexed and key:

We will store the key in sorted form in
indexTime complexity: $\Theta(\lg_2 N) + 1$ If ~~un-indexed~~ ordered and non-key.

→ Intermediate layer (Blocks of second pointers)

Time complexity $\lg_2 N + 1 + 1$

→ more than that.

e.g. Appendix of Books (secondary index)

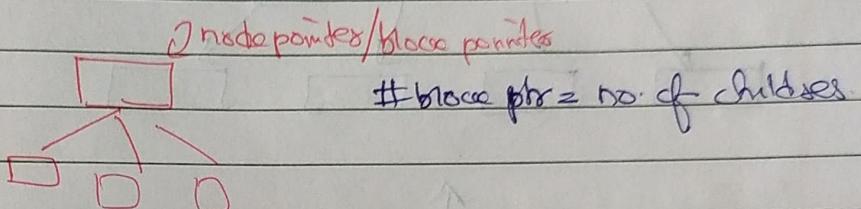
with Pg No. (Name of pointer)

"DR" 216, 222, 282

Introduction to B-tree and its Structure
(Dynamic Multilevel Index)

We have used B-tree, to perform insertion of data and indices as well in an efficient manner.

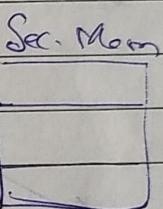
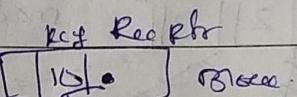
Block pointer:



Keys:

Data pointers

Record pointers

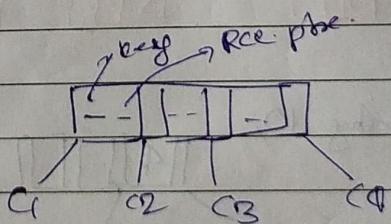


Order of a B-Tree = max no. of children a node can have

P=order of B-Tree

Children	Root	Intermediate
Max	P	P
Min	2	$\lceil P/2 \rceil$

Keys = Rec. Pts = P-1



Date 08-08-21

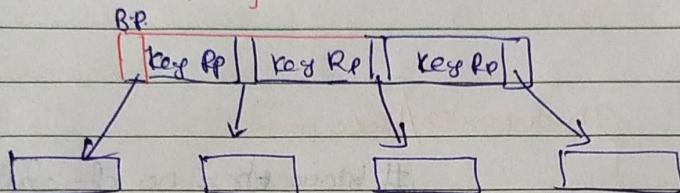
Ex:

Order = 4

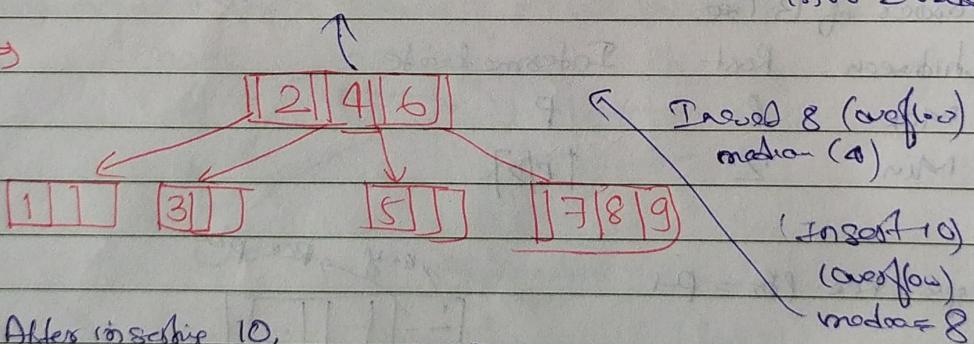
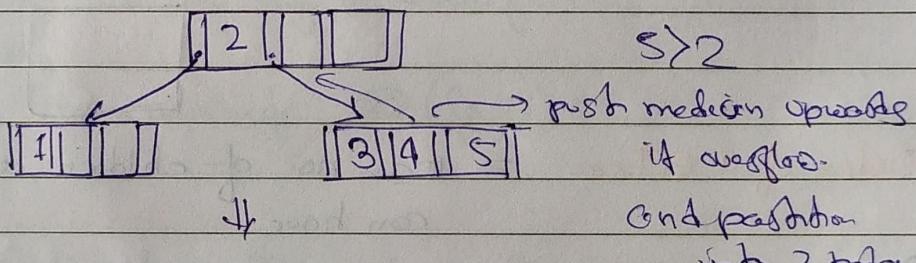
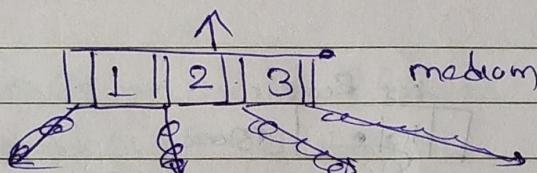
Insert values $\rightarrow 1, 2, 3, 4, 5, 6, 7, 8, 9, 10$

$$\text{Ans- Max-key} = 4 - 1 = 3$$

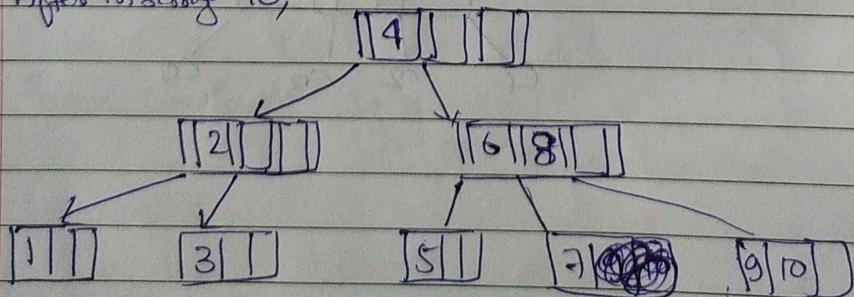
$$\text{Min-key} = \lceil \frac{9_2}{4} \rceil - 1 = 1$$



Soln.



After inserting 10,



Find order of the B-Tree

Q1 key-size of 10 B-ptr

Block size = 512 bytes, data ptr = 8 bytes

Block ptr = 5 bytes

Sol: we know,

Let $p = \text{order of tree}$

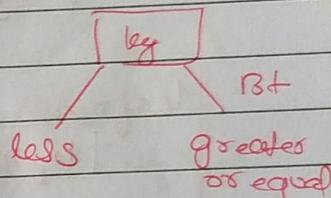
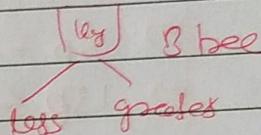
$$5 \times p + (\underbrace{(p-1) \times 8}_{\text{Block}} + \underbrace{(p-1) \times 10}_{\text{key}}) \leq 512$$

$$23p \leq 530$$

$$p \leq \frac{530}{23} \approx 23.04$$

$$(p=23)$$

24th won't fit



Difference b/w B-tree and B+ Tree

B-tree

1) Data is stored in leaf as well as internal nodes.

2) Searching is slower, deletion complex.

3) Leaf and Internal together

4) No redundant keys in B-tree

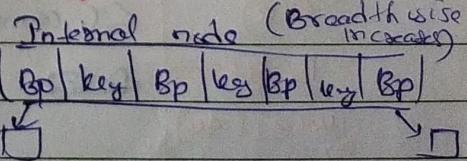
B+ Tree

1) Data is stored only in leaf nodes.

2) Searching is faster, deletion is easy (directly from leaf node).

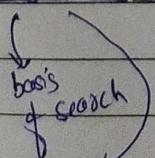
3) Leaf linked together like linked list.

4) Redundant keys may be present

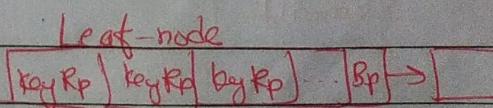


structure same for all

Bp	Key	Rp	Bp	KRp	Bp
----	-----	----	----	-----	----



p/8
to data in SM



points to next node

Order of B+ Trees

key size = 10 bytes

B size = 512 bytes

Data = 8 bytes ; block phr = 5 bytes

Find orders of leaf and non-leaf node?

Solu: Non-leaf nodes

$$B_p \times P + \text{Key}(P) \leq 512$$

$$\Rightarrow SP + 10(P) \leq 512$$

$$\Rightarrow 18P \leq 522$$

$$P \leq 28.9 \Rightarrow \boxed{P=28} \quad \text{order} = \text{no. of max children}$$

leaf nodes = max value of (key, bp) pair:

$$\boxed{\text{key} \oplus \text{by bp}} - \boxed{\text{bp}}$$

$$P(\text{key} + D_{ps}) + B_p \leq 512$$

$$\Rightarrow P(18) + S \leq 512$$

$$18P \leq 507$$

$$P < \frac{507}{18} \Rightarrow P = 28.2 \Rightarrow (P=28) \text{ order}$$

Immediate DB modification (log based recovery)

$A = 100$	T_1	logs
$B = 200$	$R(A)$	$\langle T_1, \text{start} \rangle$
	$A = f(A)$	$\langle T_1, A, 100, 200 \rangle$
	$w(B)(200)$	$\langle T_1, B, 200, 900 \rangle$
	$R(B)$	$\langle T_1, \text{commit} \rangle$
	$B = B + 200$	start
	$w(B)(400)$	
	commit	

Immediately reflect the changes in RMP to the DB

why B-tree is chosen?

If failure occurs:

- ① → if start and commit both there
→ Do REDO save all data to DB.

- ② → If commit not done

→ UNDO (save old values to DB, from logs)

* It is also known as UNDO/REDO Strategy.

View of total database content is Conceptual view

Database Supports

- (A) OLAP (B) OLTP
(C) OLAP and OLTP (D) operational DB

Which of the following does not comes under five V's of Big Data

- (a) volume (b) variety (c) velocity (d) visualization
(removes value & veracity)

Hadoop is framework that works with variety of selected tools (commonly GRP include)

- (A) MapReduce, file and Hbase

Which does not describes Hadoop

- (a) open Source (b) Real time (c) Java based
(it is batch processing)
(d) Distributed computing approach

(late)

Deferred database modification (log based recovery)

→ log is a small file which stores all actions performed.
 Log

T_i

R(A)

<T_i, S_{TiA}>

A>A100

<T_i, A1200>

W(A)

<T_i, B1900> New value

R(B)

<T_i, commit>

B>B1200

W(B)

Commit

(Saved after commit)

Recovery

If commit done :

→ RDO (Save all new values in DB)

If not committed :

→ Roll back to old values

⇒

→ Also known as NO UNDO / REDO ~~over~~ method.

Like command in SQL

→ starting from 'A'

→ where name Like 'A%' ;

→ ending with '%'

→ where name like '%n' ;

→ Contains '%'

→ where name like '%key%' ;

→ contains 'a' in 2nd place

→ where name like 'a%';

→ name contains 'a' in 2nd place and length should be 5 characters

→ where name like '_a_____';

PL SQL

① Find the Sum of two numbers

declare

a int;

b int;

c int;

begin

a := 8a; // user input

b := 8b;

c := a + b;

dbms_output.put_line ('Sum of a and b = ' || c);

end;

② Find greatest of two numbers

declare

a int;

b int;

begin

a := 8a;

b := 8b;

if (a > b)

then

dbms_output.put_line ('a is greater');

else

db

endif;

end;

For loop

Declare

a number(2);

begin

for a in 0..10

loop

dbms_output.putline(a);

end loop;

end;

While loop;

{}

while a < b

loop

~~declare~~ a:=a+1

end loop;

end;

Set server output on

Single row functions and multi row functions

(1) Single row function: Applied to single row

and gives single o/p.

(a) Number func: round, trunc, mod

(b) character func: A. Case manipulation

Upper, lower, initcap

B. Character manip.

concat, substr, length, instr, LPAD, RPAD, Trim, Replace

(friending word)

(C) Date functions

→ Months between, Add months, Next Day, Last Day,
Round, Trunc

(D) General functions (Dealing with NULL Values)

NVL, NVL2, NULLIF, COALESCE, CASE, DECODE
(Switch)

(E) Data type conversion

→ TO_CHAR, TO_NUMBER, TO_DATE

(G) Multirow functions: Applied to multiple rows but gives single dp.

e.g. Aggregate function.

Character functions

(a) Case manip.

(b) character manip

1. Concat ('Varun', 'Singh')

2. Substr ('Varun', 2, 4) ⇒ aru

3. Instr ('Varun', 'u') ⇒ 4

4. length

5. LPAD ('Varun', 10, '*') ⇒ ****Varun

6. RPAD ↑ Total length ↓ Left pad

7. Trim ('V' from 'Varun') ⇒ arun

8. Replace ('Varun', 'V', 'T') ⇒ Tarun

Abdul Wasi

B & BT Trees

Disk Structure

How data is stored on Disk

What is indexing?

What is multilevel indexing?

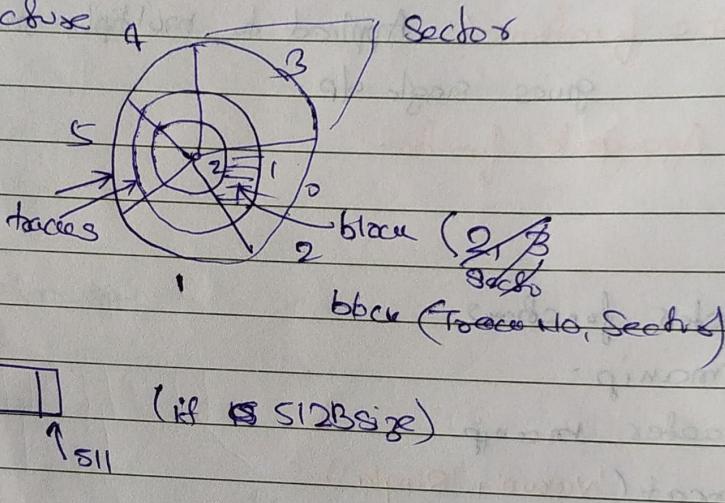
m-Way Search Trees

B-trees

Insertion & Deletion in B-trees

BT Trees

Disk Structure



DBMS: Organizing the data
in the HD, so that it
can be accessed efficiently

Data Structure: Organizing the
data in MM, that is
directly used by the program.

How data is stored in Disks?

→ Data is stored in Blocks

→ No. of tuples stored in 1 block = $\frac{\text{Block Size}}{\text{Record Size}}$

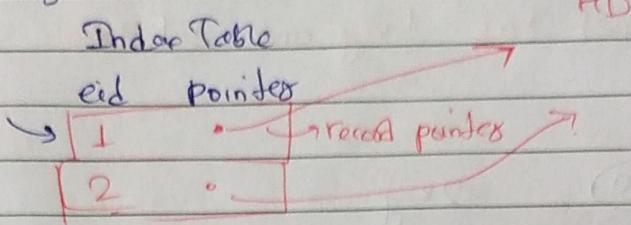
$$\text{No. of blocks req'd} = \frac{\text{Total No. of Records}}{\text{No. of records in 1 block}} = \frac{128}{128} = 4 \text{ Nos}$$

$$= \frac{100}{4} = 25 \text{ Blocks}$$

H Indexing

particulars

- # Reducing the time to access the blocks is done through INDEXING.



→ Indices are stored in Hard Disk.

(i) In Dense, No. of records in indices = # of records

Total No. of indices in 1 Block = $\frac{512}{16}$ (block size)
= 32 (index size)

$$\text{No. of indices} = 100$$

$$\text{No. of blocks needed} = \frac{100}{32} \quad \begin{matrix} \text{Total indices} \\ \# \text{ indices in 1 block} \end{matrix}$$

$$= 3.2 \approx 4 \text{ blocks}$$

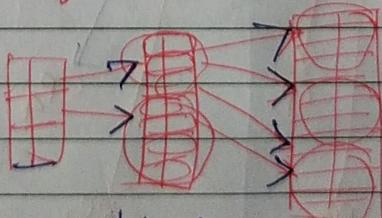
$$\text{Sq Time} = (4+1)$$

finding block

Multiple Multilevel Indexing

If searching in an index table also takes a large time (if no. of indices is huge).

This is the basis for B-tree usefulness in DBMS



Removal and
addition of
high level
indices based

on no. of entries
is done automatically.

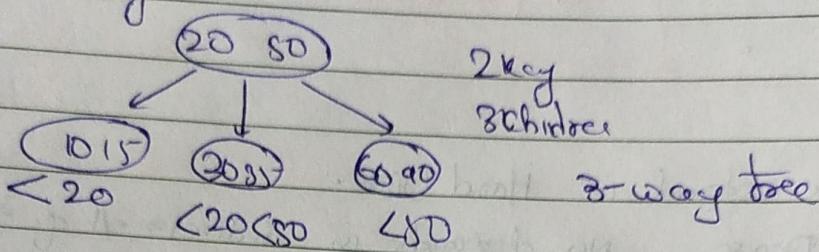
For self-managing high level
indexing, we use B-tree.

m-Way Search tree

→ Binary Search tree has 1 key and 2 children

Ex. m-Way

e.g.

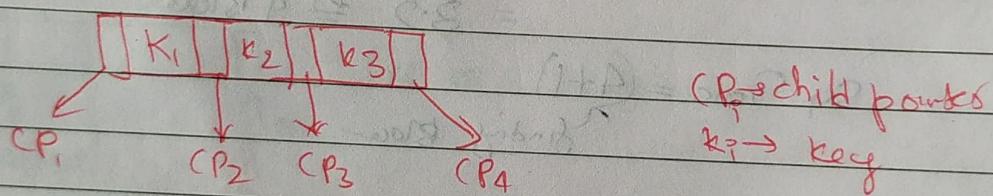


In a m-way tree

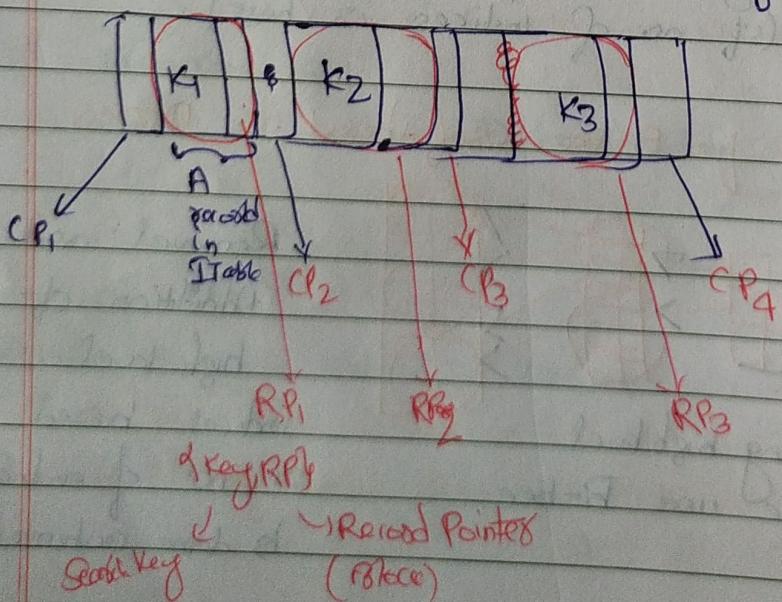
→ (m-1) keys and m-children are fixed

So Binary Search Tree is a subset of m-way tree with m=2.

Node-structure for 4-way tree

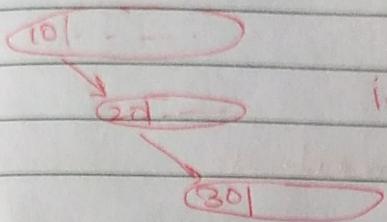


Using m-way tree for indexing.



Problem

→ There is no control of the creation of m-Way Search Tree.
For 10-way Search tree if keys are 10, 20, 30



i.e. → In worst case, for n keys height will become n , that will result into linear search.

B-trees (m-way ST + Rules)

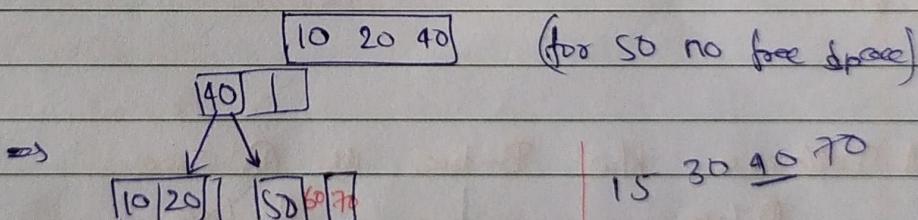
The guidelines are provided by B-Tree.

Rules:

children

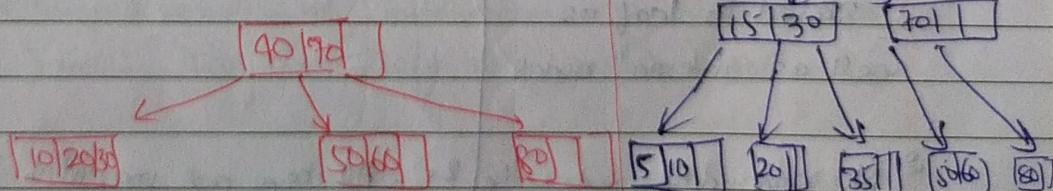
1. Every node must have $\lceil m/2 \rceil$ ~~nodes~~ children, then only child nodes are filled
2. Root can have minm ~~root~~ 2 children (1 key)
3. All leaf at same level.
4. Bottom-up creation process

e.g. B-tree! $m=4$; keys: 10, 20, 40, 50



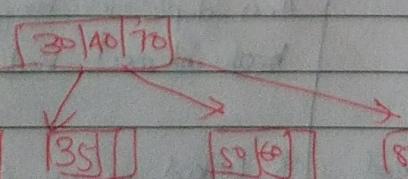
Insert 80.

↳



↳ 35

Insert 50
50 10 20 30 (biased)



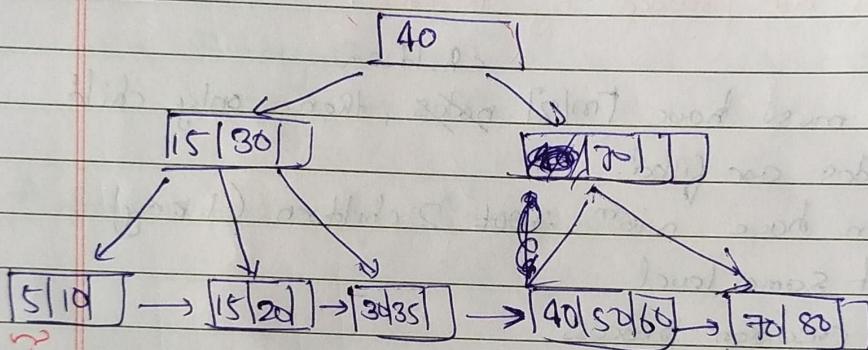
The tree grows upwards.

→ Every key or B-tree node has
(key, Rec. pointer)

B+ Tree

- Record pointers will only be present at leaf nodes
- So, every node will be present in its child leaf node as well

Prev. example :-



consist
of key and
rec. ptr.

Differences b/w B-tree vs B+ Tree

B Tree

1. Keys and records are stored in leaf as well as internal nodes.
2. No redundant node value are present.
3. Leaf nodes are not linked to each other.

B+ Tree

1. Only in leaf nodes.
2. All keys are present in leaf, their duplicates may be present in high internal nodes.
3. Linked in a sequential manner.

B-tree

BTree

4. Deletion in Btree is very complex as we need to take care of the child nodes.

4. Deletion is very simple as all keys are present in leaf nodes.

Functional dependency: is a constraint that specifies relationships b/w two sets of attributes whose one set can accurately determine another set.
 Determinant \rightarrow Dependent

E-R Model \rightarrow High level data model
 \rightarrow Conceptual design for the database
 Develops a simple and easy view of data.
 \rightarrow Pictorial Representations

Concurrency control protocols

\rightarrow Lock Based

- (i) Shared exclusive \ominus
- (ii) 2PL $\xleftarrow{\text{Granularity}} \xrightarrow{\text{Sharing}}$
- (iii) Strict 2PL
- (iv) Rigorous 2PL
- (v) Conservative 2PL

\rightarrow Time Stamp ordering ~~Order~~ protocol.

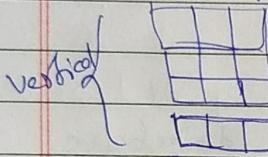
Inner Joins \Rightarrow results in intersection of two tables.
 Outer Joins: Full outer, left outer, right outer

Date: 29/08/21

Scaling: We either compress or expand the system to meet the expected needs.

Scaling operation can be achieved by adding resources to meet the smaller expectation in the current system, or by adding a new system in the existing one, or both.

Vertical Scaling: When new resources are added to the existing system to meet the expectation, it is known as vertical scaling.

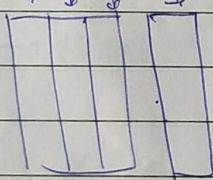


Easy and cheaper than

horizontal scaling

Horizontal Scaling

: When new server nodes are added to meet requirements.



Difficult and costlier than
vertical scaling

Sharding: Helps the system to keep data into different resource according to sharding process.

→ "shards" means "small part of a whole".

It is a type of DB partitioning which makes it more manageable and can be accessed faster.

Features of Sharding:

1. Makes the DB smaller
2. Makes the DB faster
3. DB becomes much more easily manageable
4. Can be a complex operation Sometimes
5. Reduces the transaction cost of the DB.

Denormalization: DB optimization technique, in which we add redundant data to one or more tables helps us in getting rid of costly join operations. If tables are large, we may spend an unnecessary long time in performing joins.

Under denormalization, we are okay with some redundancy and some effort to update the DB in order to get the efficiency advantages of fewer joins.

Pros:

- Retrieval time is faster
- Less bugs because of simple query retrieval.

(Cons)

- Insert and update may become costly
- Data may be inconsistent.
- Data redundancies necessitates more storage.

Intension and Extension in a DB:

Intension is the permanent part of the relation, comprises of two things relation and relation schema and the constraints Integrity constraints defining the key constraints, referential constraints.

Extension: is the snapshot of the system at a particular time. Displays values for tuples in relation at the particular instance of time. keeps changing with time as tuples are added deleted, edited.

SQL

1. Rdbms
2. Has predefined schema
3. Can require upfront preparation
4. SQL dbs are vertically scalable
(by increasing resources)
5. SQL databases are table-based
(suitable for multirow transactions)
6. Follows ACID properties

NoSQL

1. Non-Rdbms
2. Flexible schema
3. Changes can be made at any time
4. Are horizontally scalable.

5. Key-value pairs, document based

6. Breweer CAP theorem.

(Consistency, Availability, Partition tolerance)

7. Support from vendors
e.g. MySQL,
PostgreSQL

7. Only community supported
e.g. MongoDB, cassandra.

8. Not suited.

8. Best suited for hierarchical data storage.

CAP

Consistency: All client programs who are reading data from the cluster see the same data at any given pt. of time.

Availability: If we make a request, we must get a response no matter what happened inside the cluster.

Partition tolerance: Cluster (as a whole) continues to function even if there is a partition b/w two nodes.